

Advanced OS Report

倪远东 2017210845

Virtual machine monitors: Current technology and future trends

- 1960s: 硬件资源昂贵
 - 硬件复用, 诞生VMM技术
- 1980s: 成本降低、多任务操作系统
 - VMM逐渐被淡忘
- Now: 硬件成本降低, 主机服务多易崩溃, 管理困难
 - 在单台机器(或集群)上运行多个虚拟机, 每个虚拟机提供特定服务
 - 资本涌入

硬件复用 -> 提供安全性、稳定性、便捷性保障

CPU虚拟化

- 概述
 - VMM技术使得虚拟机指令在CPU上可以“直接执行”
- 挑战
 - 2005年, 没有CPU支持硬件虚拟化
- 技术
 - 半虚拟化, 将“直接执行”与“二进制翻译”相结合
- 未来趋势
 - AMD和Intel相继推出硬件虚拟化支持

内存虚拟化

- 概述
 - VMM维护一个影子内存, 用于映射虚拟机中的内存数据
- 挑战
 - VMM没有足够的信息来管理活跃分页
- 技术
 - GuestOS中运行Ballooning进程, 弹性分配
- 未来趋势
 - Extended Page Table

I/O虚拟化

- 软件模拟
 - 有效复用, 传输效率低
- 半虚拟化
 - 共享一块内存区域
- 直接分配
 - 最大化性能

总结

- Now
 - 安全性
 - 隔离性
 - 方便快捷
- Futrue
 - 应用分发的全新思路

My VM is Lighter (and Safer)
than your Container

- 虚拟机
 - 稳定性高
 - 安全
 - 镜像大、速度慢
- 容器技术
 - 资源占用低
 - 启动速度快
 - 隔离性不足

完美的LightVM在以下几点可以和容器媲美：

- 快速启动：容器可以毫秒级启动
- 高实例密度：一台主机上的容器数量级可以上万
- 暂停、恢复功能：快速暂停恢复

为什么虚拟机相对于容器的性能如此低？

- 完整的系统, 镜像比较大
- 运行的进程比较多

现象: 用于应用分发, 通常运行单个服务

改进: 精简虚拟机中的软件, 让其正好满足一个程序的运行环境

测试Xen, 找到瓶颈, 进行改进




Figure 3: The Xen architecture including toolstack, the XenStore, software switch and split drivers between the driver domain (Dom0) and the guests (DomUs).




Figure 5: Breakdown of the VM creation overheads shows that the main contributors are interactions with the XenStore and the creation of virtual devices.

测试结果

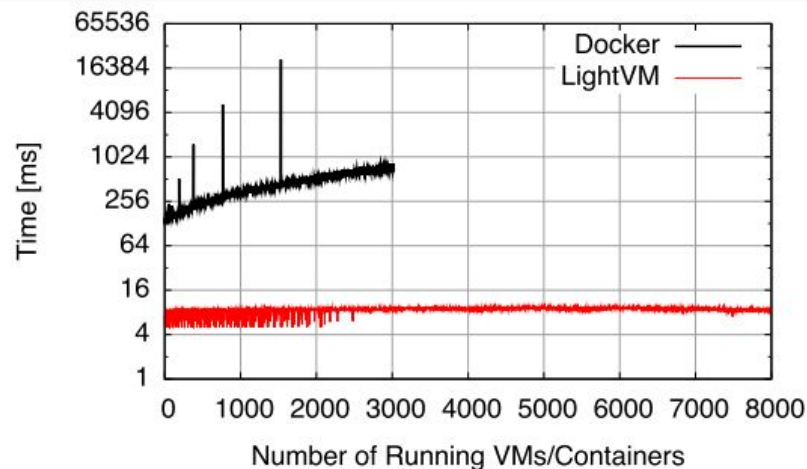


Figure 10: LightVM boot times on a 64-core machine versus Docker containers.




Figure 11: Boot times for unikernel and Tinyx guests versus Docker containers.

Making Smart Contracts Smarter

- 比特币
 - 区块链诞生
- 以太坊
 - 区块链2.0
 - 智能合约

- 基于以太坊的智能合约
 - 栈式虚拟机
 - 三种数据访问方式
 - 运行时栈
 - 内存
 - 长期存储

- 交易顺序依赖
- 时间戳依赖
- 异常处理缺失
- 重入攻击

智能合约漏洞检测工具Oyente, 采用符号执行进行分析

- CFGBuilder: 为目标合约代码构建一个较为基础CFG图
- Explorer: 遍历CFG, 执行指令
- CoreAnalysis: 漏洞分析核心模块
- Validator: 消除误报

Oyente实验

硬件: Intel i7 8550U/16GB

RAM

操作系统: Ubuntu 18.04

语言: Python 2.7.14

- 搭建环境
- 安装依赖
- 下载合约文件
- 测试本地合约
- 测试远程合约

测试本地合约

```
→ oyente git:(master) X python oyente.py -s ./source_code/DAO/DAO.sol
WARNING:root:You are using evm version 1.8.6. The supported version is 1.7.3 including Whitepaper
WARNING:root:You are using solc version 0.4.23, The latest supported version is 0.4.19
INFO:root:contract ./source_code/DAO/Token.sol:Token:
INFO:symExec: ===== Results =====
INFO:symExec: EVM Code Coverage: 56.7%
INFO:symExec: Integer Underflow: True
INFO:symExec: Integer Overflow: False
INFO:symExec: Parity Multisig Bug 2: False
INFO:symExec: Callstack Depth Attack Vulnerability: False
INFO:symExec: Transaction-Ordering Dependence (TOD): False
INFO:symExec: Timestamp Dependency: False
INFO:symExec: Re-Entrancy Vulnerability: False
INFO:symExec: ./source_code/DAO/Token.sol:44:5: Warning: Integer Underflow.
string public symbol
./source_code/DAO/Token.sol:43:5: Warning: Integer Underflow.
string public name
INFO:symExec: ===== Analysis Completed =====
INFO:root:contract ./source_code/DAO/TokenCreation.sol:TokenCreation:
INFO:symExec: ===== Results =====
INFO:symExec: EVM Code Coverage: 59.5%
INFO:symExec: Integer Underflow: True
INFO:symExec: Integer Overflow: True
INFO:symExec: Parity Multisig Bug 2: False
INFO:symExec: Callstack Depth Attack Vulnerability: False
INFO:symExec: Transaction-Ordering Dependence (TOD): False
INFO:symExec: Timestamp Dependency: False
INFO:symExec: Re-Entrancy Vulnerability: False
INFO:symExec: ./source_code/DAO/TokenCreation.sol:42:48: Warning: Integer Underflow.
/// @return Whether the token creation was successful
^
Spanning multiple lines.
./source_code/DAO/TokenCreation.sol:42:24: Warning: Integer Underflow.
/// @return Whether the token creatio
INFO:symExec: ./source_code/DAO/TokenCreation.sol:61:13: Warning: Integer Overflow.
balances[_tokenHolder] += msg.value
Integer Overflow occurs if:
balances[_tokenHolder] = 73605700760779997811274968358752567109166944192237255361232860284186811899209
./source_code/DAO/TokenCreation.sol:62:13: Warning: Integer Overflow.
totalSupply += msg.value
Integer Overflow occurs if:
totalSupply = 73605700760779997811274968358752567109166944192237255361232860284186811899209
INFO:symExec: ===== Analysis Completed =====
→ oyente git:(master) X
```

测试远程合约

```
-> oyente git:(master) X python oyente.py -ru https://gist.githubusercontent.com/loiluu/d0eb34d473e421df12b38c12a7423a61/raw/2415b3fb782f5d286777e0bcebc57812ce3786da/puzzle.sol
WARNING:root:You are using evm version 1.8.6. The supported version is 1.7.3
WARNING:root:You are using solc version 0.4.23, The latest supported version is 0.4.19
INFO:root:contract remote_contract.sol:Puzzle: Time: 03fc25909d ~ New pull request
INFO:symExec: ===== Results =====
INFO:symExec: EVM-Code Coverage: 51.0%
INFO:symExec: Integer Underflow: True
INFO:symExec: Integer Overflow: False
INFO:symExec: Parity Multisig Bug 2: False
INFO:symExec: Callstack Depth Attack Vulnerability: True
INFO:symExec: Transaction-Ordering Dependence (TOD): False
INFO:symExec: Timestamp Dependency: False
INFO:symExec: Re-Entrancy Vulnerability: False
INFO:symExec:remote_contract.sol:7:2: Warning: Integer_Underflow.
bytes public solution
INFO:symExec:remote_contract.sol:27:6: Warning: Callstack Depth Attack Vulnerability.
msg.sender.send(reward)
remote_contract.sol:20:4: Warning: Callstack Depth Attack Vulnerability.
owner.send(reward)
INFO:symExec: ===== Analysis Completed =====
```

Latest commit 03fc259 on 14 Jun 2017

Add files via upload a year ago

Add files via upload a year ago

Add files via upload a year ago

Initial commit a year ago

Add files via upload a year ago

Add files via upload a year ago

Add files via upload a year ago

README.md

- **oyente.py**: 入口函数, 可以接收以下类型输入
 - solidity合约代码
 - evm字节码
 - 远程合约代码
 - **部署在链上的合约abi**
- **symExec.py**: 执行分析的函数入口
 - build_cfg_and_analyze()
 - collect_vertices()、construct_bb()
 - full_sym_exec()
 - sym_exec_ins()

Precise and Scalable Detection of Double-Fetch Bugs in OS Kernels

- 地址空间划分
 - 用户空间
 - 内核空间
- Multi-read 的普遍性
 - 依赖查找
 - 协议/签名校验
 - 信息猜测

What is double-fetch bugs ?

```
1 static int perf_copy_attr_simplified
2 (struct perf_event_attr __user *uattr,
3  struct perf_event_attr *attr) {
4
5     u32 size;
6
7     // first fetch
8     if (get_user(size, &uattr->size))
9         return -EFAULT;
10
11    // sanity checks
12    if (size > PAGE_SIZE ||
13        size < PERF_ATTR_SIZE_VER0)
14        return -EINVAL;
15
16    // second fetch
17    if (copy_from_user(attr, uattr, size))
18        return -EFAULT;
19
20    .....
21 }
22 // Example: if attr->size is used later
23 // BUG: attr->size can be very large
24 memcpy(buf, attr, attr->size);
```

(a) C source code

Overview

- 对double-fetch进行了正式和精确的定义
- 设计工具DEADLINE来自动化发现double-fetch
- 发现了Linux kernal中的23个新bug, 以及FreeBSD中的1个新bug

形式化定义

1. 至少有两次对用户空间的读操作(multi-reads)
2. 两次读取内容必须有重叠部分(overlapped-fetch)
3. 两次读取数据必须有一定依赖(data/control)
4. 两次数据不能被证明一定是相同的


Algorithm 1: High-level procedure for *double-fetch bug* detection

In : *Kernel* - The kernel to be checked

Out : *Bugs* - The set of *double-fetch bugs* found

```
1 Bugs  $\leftarrow \emptyset$ 
2 Setf  $\leftarrow$  Collect_Fetches(Kernel);
3 for F  $\in$  Setf do
4   | Setmr  $\leftarrow$  Collect_Multi_Reads(F)
5   | for  $\langle F_0, F_1, F_n \rangle \in$  Setmr do
6   |   | Paths  $\leftarrow$  Construct_Execution_Paths(F0, F1, Fn)
7   |   | for P  $\in$  Paths do
8   |   |   | if Symbolic_Checking(P, F0, F1) == UNSAFE then
9   |   |   |   | Bugs.add( $\langle F_0, F_1 \rangle$ )
10  |   |   | end
11  |   | end
12  | end
13 end
```

从multi-reads到double-fetch



(b) Memory access patterns

```
1 // init root SR
2 $0 = $PARM(0),    @0 = $UMEM(0) // uattr
3 $1 = $PARM(1),    @1 = $KMEM(0) // attr
4 ---
5 // first fetch
6 fetch(F1) is {A = $0 + 4, S = 4}
7 $2 = @0(4, 7, U0), @2 = nil      // size
8 ---
9 // sanity checks
10 assert $2 <= PAGE_SIZE
11 assert $2 >= PERF_ATTR_SIZE_VER0
12 ---
13 // second fetch
14 fetch(F2) is {A = $0, S = $2}
15 @1(0, $2 - 1, K) = @0(0, $2, U1)
16 ---
17 // check fetch overlap
18 assert F2.A <= F1.A < F2.A + F2.S
19      OR F1.A <= F2.A < F1.A + F1.S
20 // --> satisfiable with @0(4, 7, U)
21
22 // check double-fetch bug
23 prove @0(4, 7, U0) == @0(4, 7, U1)
24 // --> fail, no constraints on @0(4, 7, U1)
```

(c) Symbolic representation and checking

Hyperkernel: Push-Button Verification of an OS Kernel

- 内核的“正确性”尤为重要
- 如何证明内核设计的“正确性”
 - 形式化验证
- 如何自动化验证, 提高可靠性和效率
 - Hyperkernal

设计这样一个自动化工具面临着三大挑战

- 设计需要在可用性和自动化证明之间进行折中
- 分离内存中的用户空间与内核空间地址，简化证明中的内存管理
- C语言会使得推理复杂，如何建模

如何设计有限的接口：

- 强制资源被引用的生命周期
- 强制进行细粒度的保护
- 验证链接资源的结构

例：FD机制返回最小可用文件描述符

用状态机来描述期望的内核行为：

- 抽象的内核状态规范
- 声明的规范

给出了两个定理, 使用Z3进行定理验证

Hyperkernal实验

硬件: Intel i7 8550U/16GB

RAM

操作系统: Ubuntu 18.04

- 搭建环境
- 安装依赖
- *make*
- *make qemu*
- *make hv6-verify*

```
C++      irpy/compiler/irpy
IRPY     o.x86_64/hv6/hv6.py
Parsing took 37.945 ms.
terminate called after throwing an instance of 'std::invalid_argument'
what(): stoi
Aborted (core dumped)
Makefile:208: recipe for target 'o.x86_64/hv6/hv6.py' failed
make: *** [o.x86_64/hv6/hv6.py] Error 134
```

The Scalable Commutativity Rule: Designing Scalable Software for Multicore Processors

如何评价软件的可扩展性？

- 软件架构设计
- **接口设计？**

如果一组操作在内存访问上是没有冲突的, 我们就认为这组操作时可扩展的

- SIM交换性规则
- 可扩展性规则

设计可扩展接口的策略：

- 分解复合的操作
- 去掉接口中的一些并不需要的功能
- 弱化不必要的排序约束
- 异步释放资源

- ANALYZER:自动分析接口的复杂性, 给出交换条件
- TESTGEN:用于从ANALYZER给出的交换条件中生成大量的测试用例
- MTRACE:将测试用例运行在实际的系统上, 并检查实现是否无冲突